


Comprehension and reachable jailbreak vary across LLMs — yet the winning attack lever is consistent within each family

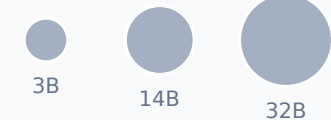
Lingua-SafetyBench: one marker per model, area $\propto$ parameters; shaded region = model family



Consistent lever per family

- Gemma** → persona
- Qwen** → role separation
- Llama** → multilingual / amount

Marker area $\propto$ parameters



⇒ Attack must be per-target & comprehension-aware.