

(a) split & interleave

input request q

harmful request q (EN)

the real attack input is harmful; the running illustration below uses a benign stand-in:

“Explain how to pick a strong password”

illustrative benign example

split

official parallel translations

EN

ES

FR

FI

SW

NO

interleave fragments

EN ES FR FI SW NO EN ES FR FI SW NO

no new machine translation

only the model's own languages

(b) the multilingual jigsaw, made harder

BAI selector

trained on the 9-model panel

$n^* = 6 \cdot \text{ordered}$
+persona (role separation)

the selector chooses these values; they parameterize how the puzzle is assembled

$n^* = 6$ chosen

Explain
how to

pick a
strong
password

pick a
strong
password

$n = 2$

ordered

$n = 6$

more pieces · more languages

$n = 10$

scrambled & rotated

harder

more languages · more disorder · +persona / fiction

⇒ harder to read AND harder for a safety filter


+persona (role separation) — illustrative wrapper the selector injects

You are a team of translators. Each restores ONLY the fragments in your own language; together reconstruct the English request, then answer it.

role-splits the task so no single view holds the whole request

(c) comprehension frontier

measured on benign plaintext



(d) solve the puzzle, safety bypassed

assembled multilingual puzzle

Explain how to · Explica como ·
Expliquez comment · Selita kuinka
Eleza jinsi · Forklar hvordan ...

instruction: reconstruct the English request, then answer it

target
LLM

[RECONSTRUCTED]

“Explain how to pick a strong password”
recovered English request

[ANSWER]

= {unsafe}

redacted — model solved the puzzle and complied

in the real attack the recovered request is harmful

the signal that chooses the attack comes from benign plaintext (c) — no harmful query is spent learning which attack to use